


Gradient variance of the deployed CVaR objective versus qubit count
 exact parameter-shift gradient; $\theta \sim \mathcal{N}(0, 0.6^2)$; deployed rank-ladder H; $T = 0$
 error bars = bootstrap 95% CI over the θ draws; slopes = least squares on $\log_2$ Var, with CI



- $\alpha = 1.0$ (linear cost, $\langle \psi | H | \psi \rangle$); slope $-0.64 \log_2/\text{qubit}$ $[-0.71, -0.59]$
- $\alpha = 0.5$; slope $-0.60 \log_2/\text{qubit}$ $[-0.67, -0.54]$
- $\alpha = 0.25$; slope $-0.30 \log_2/\text{qubit}$ $[-0.37, -0.23]$
- $\alpha = 0.1$; slope $-0.08 \log_2/\text{qubit}$ $[-0.16, +0.24]$